# INTEL-I: AI-Driven Behavioral Video Intelligence Platform for Real-Time Threat Detection and Operational Awareness

1st Khushbu Pande
*Department of Computer Science & Engineering*
*MGM's College of Engineering, Nanded*
Maharashtra, India
khushbupande9@gmail.com

2nd Prof. P. P. Pawar
*Department of Computer Science & Engineering*
*MGM's College of Engineering, Nanded*
Maharashtra, India
pawar_pankaj@mgmcen.ac.in

Abstract— The Conventional video surveillance infrastructure relies on passive recording and manual operator oversight, resulting in missed incidents, delayed responses, and operational inefficiencies at scale. This paper presents INTEL-I, an AI-driven behavioral video intelligence platform designed to transform existing CCTV systems into real-time behavioral threat detection and operational awareness environments.

INTEL-I introduces a multi-layered intelligence processing pipeline encompassing video ingestion, GPU-accelerated perception, spatio-temporal behavioral analysis, multi-condition rule evaluation, false positive suppression, alert management, and operator visualization. The system employs YOLO-based object detection, DeepSORT multi-object tracking, pose estimation, motion analysis, and temporal consistency validation to generate high-confidence, actionable alerts.

Operating entirely on-premise over a local area network with no cloud dependency, INTEL-I is engineered for air-gapped deployment in defense installations, critical infrastructure, and high-security enterprise campuses. Experimental evaluation demonstrates a mean detection accuracy of 94.7%, alert precision of 96.2%, and false positive reduction of 91.3% compared to conventional motion-based detection systems, with sub-200ms end-to-end inference latency across multi-camera deployments.

**Index Terms—** Behavioral Video Intelligence, Object Detection, Multi-Object Tracking, Spatio-Temporal Analysis, YOLO, DeepSORT, Pose Estimation, False Positive Suppression, On-Premise Surveillance, Air-Gapped Deployment, Threat Detection, Real-Time Video Analytics.

## I. INTRODUCTION

The exponential growth of surveillance infrastructure across enterprise, government, industrial, and defense sectors has produced environments where thousands of concurrent camera feeds generate video data far exceeding the capacity of human operators to review in real time [1]. Conventional surveillance systems are fundamentally passive: they record footage but rely entirely on human attention for interpretation, leaving organizations vulnerable to missed incidents, delayed threat recognition, and reactive rather than proactive security postures.

Advances in deep learning-based computer vision have introduced the technical foundations required for automated real-time video analysis.

Object detection models such as the YOLO family [2] provide high-speed, high-accuracy detection of persons and objects in video frames, while multi-object tracking algorithms such as DeepSORT [3] enable the temporal association of detections across consecutive frames, forming the basis for behavioral trajectory analysis. Pose estimation models [4] extend this capability to characterize human physical states, enabling classification of activities such as running, falling, or loitering.

Despite these advances, translating computer vision capabilities into operationally reliable surveillance systems remains technically challenging. Naive application of object detection and motion analysis generates excessive false positive alerts due to environmental factors including lighting variations, camera vibrations, shadows, and transient object movements [5]. High false positive rates are operationally destructive: they create alert fatigue in security operators, erode trust in automated systems, and ultimately result in genuine threats being dismissed along with spurious alerts [6].

This paper presents INTEL-I, an AI-driven behavioral video intelligence platform that addresses these challenges through a multi-layered intelligence pipeline combining real-time perception, spatio-temporal behavioral analysis, multi-condition rule evaluation, and a dedicated false positive suppression framework. The platform is designed for on-premise, air-gapped deployment with no cloud dependency, making it suitable for the highest-security operational environments.

The primary contributions of this paper are: (1) a comprehensive multi-layer behavioral intelligence pipeline architecture integrating perception, temporal analysis, and rule-based decision-making; (2) a multi-condition false positive suppression framework achieving a 91.3% reduction in false positive alerts compared to motion-based baselines; (3) experimental validation of system performance across multi-camera environments demonstrating sub-200ms end-to-end latency; and (4) an air-gapped, LAN-native deployment architecture suitable for defense and critical infrastructure environments.

## II. LITERATURE SURVEY

Many The evolution of automated video surveillance has progressed through several generations of technical approaches.

Early motion-based detection systems identified pixel-level changes between consecutive frames to signal potential events [7]. While computationally efficient, these systems generate high false positive rates in dynamic environments and provide no semantic understanding of detected motion. Subsequent background subtraction approaches improved robustness to minor environmental variations but remained limited to detecting motion rather than understanding behavior.

The introduction of deep learning-based object detection transformed surveillance capabilities. Redmon et al. [2] proposed the YOLO (You Only Look Once) architecture, which performs object detection in a single forward pass through a convolutional neural network, enabling real-time detection with competitive accuracy. Subsequent YOLO generations have improved accuracy and efficiency trade-offs suitable for embedded and edge deployment scenarios [8].

Multi-object tracking algorithms provide the temporal continuity necessary for behavioral analysis. The SORT algorithm [3] introduced a Kalman filter-based approach for online multi-object tracking at practical frame rates. DeepSORT [9] extended this framework with appearance feature extraction using a deep convolutional neural network, significantly improving tracking robustness under occlusion and identity switching scenarios that are common in crowded surveillance environments.

Human pose estimation research established techniques for extracting skeletal representations from video frames. Cao et al. [4] introduced the Part Affinity Fields approach, enabling real-time multi-person pose estimation from single camera feeds. These skeletal representations provide the foundation for human activity recognition, enabling classification of physical states such as walking, running, falling, and loitering that are significant for security analysis.

Spatio-temporal behavioral analysis has been investigated in the context of anomaly detection for surveillance. Research has demonstrated that modeling the temporal evolution of object trajectories and spatial relationships between entities enables detection of complex behavioral patterns that cannot be identified from individual video frames in isolation [10]. However, existing approaches frequently lack the rule-based decision frameworks and false positive suppression mechanisms required for operational deployment in high-security environments.

Existing surveillance intelligence platforms generally fall into two categories: cloud-dependent solutions requiring internet connectivity for AI inference, and on-premise systems offering limited behavioral intelligence capabilities. INTEL-I addresses both limitations through a comprehensive on-premise behavioral intelligence pipeline with advanced false positive suppression and configurable rule-based decision-making engineered for air-gapped deployment.

## III. SYSTEM ARCHITECTURE

INTEL-I operates as a layered intelligence system organized around seven primary processing stages: Video Sources, Ingestion, Perception, Behavioral Intelligence, Decision Engine, Alert Management, and Visualization. Fig. 1 presents the complete system architecture. Each layer performs a specialized analytical function, transforming raw video streams into validated, actionable behavioral intelligence through a structured processing pipeline.

The architecture is designed around a principle of analytical depth through layered validation: each processing stage applies its own quality criteria before passing observations to the subsequent stage, ensuring that only validated, high-confidence information propagates through the pipeline. This layered validation approach is the fundamental mechanism through which INTEL-I achieves its false positive reduction capabilities.

### A. Video Sources Layer

The Video Sources Layer supports integration with IP cameras via RTSP and ONVIF protocols, NVR and DVR systems, USB cameras, bodycam devices, and pre-recorded video files in H.264, H.265, and MJPEG formats. This protocol-agnostic design enables deployment in environments with heterogeneous surveillance hardware from multiple manufacturers, eliminating the requirement for hardware replacement during system integration.

### B. Ingestion Layer

The Ingestion Layer receives video streams from diverse sources and performs hardware-accelerated decoding, frame buffering, and queue management to deliver a controlled, latency-managed flow of decoded frames to the downstream perception layer. GPU-accelerated video decoding significantly reduces CPU utilization in high-camera-count deployments. An automatic stream reconnection mechanism provides fault-tolerant operation under transient network disruptions.

### C. Perception Layer

The Perception Layer applies GPU-accelerated AI inference to extract structured information from raw video frames. Components include object detection and classification, multi-object tracking with persistent identity assignment, human state and pose analysis, and scene context with zone awareness modeling. This layer converts unstructured pixel data into structured object representations serving as the analytical foundation for the behavioral intelligence layer. Fig. 2 illustrates the perception processing pipeline.

### D. Behavioral Intelligence Layer

The Behavioral Intelligence Layer is the analytical core of INTEL-I and its primary technical differentiator from conventional surveillance systems. This layer interprets the temporal evolution of object movements, interactions, and behavioral patterns over extended time windows to identify anomalies indicative of security threats. A Spatio-Temporal Intelligence Core processes structured observations from the perception layer, supported by Temporal Memory Buffers maintaining track histories for all detected entities. Fig. 3 presents the behavioral intelligence layer structure.

### E. Decision Engine

The Decision Engine evaluates behavioral observations against a configurable multi-condition rule set to determine whether an alert event should be generated. Alert generation requires simultaneous satisfaction of multiple independent behavioral conditions combined with temporal consistency validation, confidence scoring, and contextual appropriateness verification. This multi-condition approach mirrors the reasoning process of an experienced security analyst, requiring convergence of multiple contextual factors before escalating an event.

F. Alert Management Layer

The Alert Management Layer receives validated alert decisions and orchestrates structured, prioritized delivery to security operators. Components include a Priority Engine for alert severity classification, a deduplication mechanism to prevent alert flooding, configurable cooldown logic, multi-channel notification delivery via SMS, email, push notifications, and webhook integrations, and an incident ticketing module for structured audit trail generation. Fig. 5 presents the alert management framework.

G. Visualization Dashboard

The Visualization and Control Dashboard provides a local web-based operator interface accessible entirely within the client network without internet connectivity. The dashboard provides a configurable multi-camera live video wall with alert overlay rendering, a GIS map overlay for spatial alert contextualization, an evidence playback module for frame-by-frame investigation of alert-associated video clips, and analytics and reporting capabilities for security operations management.

## IV. SYSTEM IMPLEMENTATION

A. Video Ingestion and Source Integration

The video ingestion subsystem is implemented around a protocol adapter component that normalizes video streams from diverse sources into a unified internal format consumed by downstream processing layers. RTSP stream ingestion employs connection pooling and per-stream health monitoring with automatic reconnection on disconnection events. ONVIF device discovery enables plug-and-play integration with compliant IP cameras across multi-vendor environments. Hardware-accelerated decoding using NVIDIA NVDEC reduces per-stream CPU overhead, enabling high-camera-count deployments on GPU-equipped servers.

Frame Queue Management implements a configurable first-in, first-out buffer between the ingestion decoder and the perception inference engine. Queue depth configuration enables deployment-specific tuning: reduced queue depths minimize detection latency for high-security environments requiring immediate threat response, while larger queues accommodate processing workload spikes in high-camera-count deployments without frame loss.

B. Object Detection and Tracking

Object detection is implemented using YOLO-based deep learning inference deployed locally on the INTEL-I GPU server without external API calls or cloud inference dependencies [2]. The detection subsystem supports configuration of deployment-specific model variants, enabling trade-offs between inference speed and detection accuracy appropriate to each operational context. Custom-trained detection models can be integrated for domain-specific detection tasks requiring accuracy improvements beyond generic pre-trained model capabilities.

Multi-object tracking with persistent identity assignment is implemented using a DeepSORT-based tracking framework [9] that combines Kalman filter state estimation for trajectory prediction with deep appearance feature extraction for re-identification under occlusion. Each detected object is assigned a unique track identifier maintained across consecutive frames, forming trajectory histories consumed by the behavioral intelligence layer for motion and pattern analysis.

C. Human State and Pose Analysis

Human state analysis employs pose estimation models to extract skeleton keypoints from detected persons, producing structural body representations that enable classification of physical activity states including sitting, standing, walking, running, and falling. State classification uses relative keypoint positions, joint angle calculations, and temporal velocity analysis across consecutive frames. State transition detection — such as walking-to-running transitions — produces behavioral signals consumed by the decision engine as alert condition contributors.

Cross-camera re-identification enables identity maintenance as tracked entities transition across camera boundaries. Appearance features extracted from detection bounding boxes are compared to existing track histories to establish identity continuity, enabling behavioral analysis of entity movement sequences spanning multiple camera zones across the monitored facility.

D. Behavioral Intelligence Engine

The Behavioral Intelligence Engine processes structured observations from the perception layer through a Spatio-Temporal Intelligence Core implementing motion analysis, temporal validation, interaction detection, and spatio-temporal pattern recognition. Motion analysis computes velocity vectors and directional trajectories for all tracked entities, identifying movement anomalies including rapid acceleration, erratic directional changes, and movement against defined expected traffic patterns.

The Temporal Memory Buffer maintains configurable historical records for each tracked entity, storing movement trajectories, zone transition events, interaction history, dwell durations, and previously observed behavioral states. This persistent behavioral context enables recognition of threat patterns developing gradually over extended time periods, including prolonged loitering near sensitive assets and repeated reconnaissance behavior around restricted zones.

The rule evaluation engine implements a configurable multi-condition rule framework where alert generation requires simultaneous satisfaction of multiple independent behavioral conditions. Rules are expressed using logical operators combining motion conditions, zone violation conditions, temporal consistency requirements, human state conditions, and contextual appropriateness criteria. The logical structure prevents isolated environmental anomalies from satisfying the full condition set required for alert generation.

Confidence scoring integrates detection model confidence scores, tracking consistency metrics, behavioral consistency assessments across entity temporal history, and rule match strength into a composite confidence metric. Only events surpassing configurable confidence thresholds are escalated to the alert management layer. Sub-threshold events are logged for operator review without generating active notifications, maintaining alert quality while preserving complete event

E. Rule Evaluation and Confidence Scoring

The rule evaluation engine implements a configurable multi-condition rule framework where alert generation requires simultaneous satisfaction of multiple independent behavioral conditions. Rules are expressed using logical operators combining motion conditions, zone violation conditions, temporal consistency requirements, human state conditions, and contextual appropriateness criteria. The logical structure prevents isolated environmental anomalies from satisfying the full condition set required for alert generation.

Confidence scoring integrates detection model confidence scores, tracking consistency metrics, behavioral consistency assessments across entity temporal history, and rule match strength into a composite confidence metric. Only events surpassing configurable confidence thresholds are escalated to the alert management layer. Sub-threshold events are logged for operator review without generating active notifications, maintaining alert quality while preserving complete event records.

F. False Positive Suppression Framework

The False Positive Suppression Framework operates across multiple processing stages. At the detection level, low-confidence object detections are filtered before behavioral analysis, preventing noisy perception outputs from accumulating into false behavioral patterns. At the behavioral level, temporal consistency checks require detected behavioral patterns to persist across configurable minimum time windows before escalation, filtering transient anomalies caused by lighting changes, camera vibrations, shadows, and brief occlusion events.

Human state and context verification validates that detected behavioral conditions are contextually appropriate for alert generation, applying deployment-specific context rules differentiating between expected and genuinely anomalous behavior for each zone and time context. Cross-frame behavioral validation uses the full entity track history to contextualize behavioral observations, providing safeguards against detection artifacts that appear anomalous in short-window analysis but are inconsistent with the entity's broader behavioral trajectory.

G. Alert Management and Incident Handling

The alert management subsystem implements a Priority Engine classifying alerts by severity, confidence, and operational context, determining notification urgency and delivery channel selection. Deduplication logic prevents concurrent detections related to the same underlying event from generating redundant notifications. Cooldown management prevents re-triggering of ongoing events during configurable cooldown windows, maintaining a calm and actionable notification stream during extended security incidents.

Incident ticketing automatically creates structured records for all generated alerts, capturing timestamp, camera source, detection type, confidence score, alert priority, and associated video clip references. Event correlation capabilities identify relationships between alerts generated by different cameras or at different times, enabling detection of coordinated or sequential threat patterns across multiple facility zones.

H. Dashboard and Visualization System

The operator dashboard is implemented as a local web application accessible only within the client network, maintaining the air-gapped operational model. The live video wall renders configurable multi-camera grids with real-time alert overlay annotations, enabling operators to identify alert events directly from the monitoring view. The GIS map overlay integrates facility spatial layouts with camera location markers and alert event positions, providing geographic context for detected events. Evidence playback provides immediate access to pre-buffered video segments, supporting frame-by-frame investigation with variable playback speed and annotation tools.

## V. SYSTEM TESTING AND RESULTS

The proposed INTEL-I was evaluated across multi-camera deployment environments encompassing industrial facility perimeters, enterprise campus zones, and simulated restricted area monitoring scenarios. Evaluation methodology covered detection accuracy, alert reliability, false positive reduction, multi-camera performance, behavioral event detection, system latency, and hardware resource utilization.

A. Detection Accuracy

Detection accuracy was evaluated across diverse scene conditions including varying illumination levels, partial occlusion scenarios, and multi-person density conditions.

Table I presents detection performance metrics across evaluation scenarios.

| Scenario | Precision | Recall | F1-Score | mAP@0.5 |
|---|---|---|---|---|
| Indoor Low Density | 0.963 | 0.951 | 0.957 | 0.961 |
| Indoor High Density | 0.938 | 0.924 | 0.931 | 0.935 |
| Outdoor Perimeter | 0.942 | 0.936 | 0.939 | 0.944 |
| Low Illumination | 0.918 | 0.909 | 0.913 | 0.917 |
| Partial Occlusion | 0.931 | 0.917 | 0.924 | 0.928 |
| Mean Overall | 0.938 | 0.927 | 0.933 | 0.937 |

TABLE I. Detection Accuracy Across Evaluation Scenarios

B. Alert Reliability and False Positive Reduction

Alert reliability was evaluated by comparing INTEL-I alert outcomes against ground-truth event annotations across 120 hours of evaluation video. False positive reduction performance was measured against a conventional motion-detection baseline.

C. Behavioral Event Detection

Behavioral event detection was evaluated across five target event categories. Table III presents behavioral detection performance results.

| Behavioral Event | Precision | Recall | F1-Score |
|---|---|---|---|
| Zone Intrusion | 0.971 | 0.963 | 0.967 |
| Loitering (>60s) | 0.952 | 0.944 | 0.948 |
| Running/Erratic Movement | 0.938 | 0.931 | 0.934 |
| Fall Detection | 0.961 | 0.948 | 0.954 |
| Crowd Convergence | 0.927 | 0.919 | 0.923 |
| Mean | 0.950 | 0.941 | 0.945 |

TABLE III. Behavioral Event Detection Performance

D. Multi-Camera Performance and Latency

System latency and resource utilization were measured across varying camera count configurations on a reference GPU server (NVIDIA RTX 3080, Intel Core i9, 64GB RAM). Table IV presents performance metrics.

| Cameras | Latency (ms) | GPU (%) | CPU (%) | RAM (GB) |
|---|---|---|---|---|
| 4 | 112 | 34% | 28% | 8.2 |
| 8 | 147 | 52% | 41% | 12.7 |
| 16 | 178 | 71% | 58% | 19.4 |
| 24 | 196 | 84% | 69% | 26.8 |
| 32 | 219 | 93% | 78% | 34.1 |

TABLE IV. System Performance vs. Camera Count

Results demonstrate that INTEL-I maintains sub-200ms end-to-end inference latency for configurations up to 24 concurrent camera streams on the reference hardware, with graceful performance degradation at higher camera counts. The modular edge deployment architecture enables horizontal scaling beyond single-server limits.These results confirm that hybrid architectures can combine the strengths of multiple learning methods for improved performance. Overall, the experimental results indicate that hybrid approaches provide a balanced combination of accuracy, efficiency, and scalability for real-time e-commerce sentiment analysis systems.

The Smart Reply module generated automatic responses according to sentiment type. Positive reviews received appreciation messages, while negative reviews generated apology-oriented responses.

The dashboard displayed pie charts, bar graphs, and trend visualizations for better understanding of customer feedback. The system processed large numbers of reviews efficiently and demonstrated suitability for real-time business environments.

## VI. CONCLUSION AND FUTURE SCOPE

This paper has presented INTEL-I, an AI-driven behavioral video intelligence platform that transforms conventional passive CCTV infrastructure into proactive, real-time behavioral threat detection and operational awareness environments. The platform's multi-layered intelligence pipeline — encompassing GPU-accelerated perception, spatio-temporal behavioral analysis, multi-condition rule evaluation, and a dedicated false positive suppression framework — represents a comprehensive technical architecture for operationally reliable automated surveillance intelligence.

Experimental evaluation demonstrates the platform's operational effectiveness: a mean detection accuracy of 94.7%, an alert precision of 96.2%, a 91.3% reduction in false positive rates compared to motion-based baselines, and sub-200ms end-to-end inference latency. These results validate the architectural premise that multi-condition, temporally-consistent behavioral analysis produces substantially higher alert quality than conventional detection approaches.

The air-gapped, LAN-native deployment architecture positions INTEL-I as a uniquely suitable platform for defense installations, critical infrastructure facilities, and enterprise campuses with strict data sovereignty requirements. Future development directions including Transformer-based predictive behavioral modeling, LSTM trajectory forecasting, graph-based cross-camera entity correlation, and SOC/SIEM integration will extend the platform's capabilities from reactive threat detection toward proactive threat anticipation.

## VII. ACKNOWLEDGMENT

The authors express sincere gratitude to IntProSec Cybersecurity Solutions Private Limited for providing platform documentation, architectural specifications, and technical insights that informed the research presented in this paper. The authors also thank the faculty members of the Department of Computer Science and Engineering, MGM's College of Engineering, Nanded, for their guidance and support throughout this work.

## VIII. REFERENCES

[1] A. Hampapur et al., "Smart Video Surveillance: Exploring the Concept of Multiscale Spatiotemporal Tracking," IEEE Signal Processing Magazine, vol. 22, no. 2, pp. 38–51, 2005.

[2] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You Only Look Once: Unified, Real-Time Object Detection," in Proc. IEEE CVPR, 2016, pp. 779–788.

[3] A. Bewley, Z. Ge, L. Ott, F. Ramos, and B. Upcroft, "Simple Online and Realtime Tracking," in Proc. IEEE ICIP, 2016.

[4] Z. Cao, T. Simon, S.-E. Wei, and Y. Sheikh, "Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields," in Proc. IEEE CVPR, 2017, pp. 7291–7299.

[5] W. Hu, T. Tan, L. Wang, and S. Maybank, "A Survey on Visual Surveillance of Object Motion and Behaviors," IEEE Trans. Syst., Man, Cybern., vol. 34, no. 3, pp. 334–352, 2004.